

AI-ready metadata in practice

Lessons from building Portolan

Cayetano Benavent

CARTO and Portolan contributor

Eurac Research, Bolzano / Bozen

7-8 September 2026

Who is talking

Cayetano Benavent

Director of Engineering, Professional Services, at
CARTO

 Based in Seville, Spain

 Twenty-plus years in the geospatial industry

 Long-time contributor to the open-source software community

 **Portolan** SDI project collaborator

FIND ME

email cayetano@cartodb.com

github github.com/cayetanobv

linkedin linkedin.com/in/cayetano-benavent-viñuales-8b8416a

THIS TALK

slides cayetanobv.github.io/ogc-metadata-summit-bolzano

catalog github.com/cayetanobv/south-tyrol-geodata-portolan-mirror

project portolan-sdi.org

What Portolan is

Portolan is an **open-source specification and toolkit** that combines **cloud-native file formats** with **clear metadata** and **documentation** so that **people and agents** can publish and use data **at any scale**.



specification and toolkit

spec v0.2.0 · portolan
CLI · rashid validator ·
a registry of catalogs



cloud-native file formats

GeoParquet for vector ·
COG for raster · PMTiles
for display



clear metadata

a STAC 1.1 publishing
profile: what a catalog
must contain



documentation

README.md for people ·
AGENTS.md for agents,
at every level



people and agents, any scale

ten agent skills · from a
city's open data to a 369
TB archive

Why Portolan: the same goals as your SDI, a lighter way to deliver them

SAME VISION · A FEDERATED SPATIAL DATA INFRASTRUCTURE, DISCOVERABLE, INTEROPERABLE, OPENLY ACCESSIBLE TO ALL



AI-ready

TODAY

SDIs were not built for agents. Agents need full data and rich metadata, fast.

INSTEAD

Agents are first-class readers.



Easy to implement

TODAY

Database, WFS, WMS, schemas, specialist staff. Most organisations cannot take part.

INSTEAD

Files in a bucket.



Scalable

TODAY

Agents crawl whole datasets. Popularity means more servers, until something breaks.

INSTEAD

Scales with nobody on call.



Low cost

TODAY

Servers always on, ops engineers, a traffic spike is an emergency.

INSTEAD

Storage and egress.



Sovereign

TODAY

Dependence on foreign clouds and vendors.

INSTEAD

The full stack in your jurisdiction.

Two decades of SDI work got the hard part right: models, catalogues, INSPIRE themes, ISO records. What changed is the delivery layer underneath. Your catalogues and records stay authoritative.

You already know the building blocks. Most carry an OGC stamp.

WHAT A PORTOLAN CATALOG IS MADE OF

STAC 1.1 catalog · collection
· item · asset

OGC Community Standard
Oct 2025

GeoJSON · every Item is a
Feature

IETF RFC 7946

Cloud Optimized GeoTIFF

OGC Standard
2023

GeoParquet · vector

OGC SWG
approval pending

PMTiles · map display

community format

Zarr · COPC planned

Zarr: OGC Community Standard
2022

HTTP range requests ·
object storage · CORS

IETF · W3C

Nothing here is ours.

STAC 1.1.0 BASELINE + PORTOLAN v0.2.0 PROFILE



DESCRIBES + LINKS

Catalog → Collection → Item → Asset

same STAC objects



PORTOLAN

v0.2.0 · pre-1.0

PROFILES THE PUBLICATION

opinionated · versioned · checkable

THE PROFILE MAKES SIX PUBLISHING CHOICES EXPLICIT

● STRUCTURE + LINKS

● FORMATS + STATISTICS

● ACCESS + HOSTING

● LICENSE + PROVENANCE

● README + AGENTS


● VISUALIZATION

schemas.portolan-sdi.org/portolan/v0.2.0/schema.json

"STAC describes the data. Portolan defines how it is published." A publishing profile on top of these standards, not a new format.

Portolan connects the pieces: publish · update · use

ONE PRACTICE ACROSS THE DATA LIFECYCLE



PORTOLAN

STAC 1.1 · GeoParquet · COG · PMTiles · HTTP · object storage
opinionated · checkable · no new data format · nothing replaced

AI-READY

EASY TO
IMPLEMENT

SCALABLE

LOW COST

SOVEREIGN

What Portolan adds: four things the standards leave open, made mandatory



A quality floor you can check

bbox columns and spatial order in every GeoParquet ·
overviews in every COG · providers and licence on every
collection · range requests and CORS on the host. `rashid`
reads the bytes, not just the JSON.



Two documents at every level

`README.md` for the person deciding whether to trust the
data · `AGENTS.md` for the agent that needs the join key and
the caveats.



No server in the request path

a folder in storage the publisher controls · cost is storage
and egress · nothing runs after publishing.



Federation without a central copy

a registry of independently hosted catalogs, plain JSON in
version control · producer, provider and host named on
each.

DOES A STATIC MODEL SCALE? FIELDS OF THE WORLD, ONE PORTOLAN CATALOG ON SOURCE COOPERATIVE

369 TB

published

106.6 TB

served in the 28 days to 27 August 2026

0

servers in front of it

Data: Fields of the World (fieldsoftheworld.org), published on Source Cooperative; figures from the Portolan launch post, 2 Sep 2026.

Story 1 · A mirror of this province's open data, published as a spec v0.2.0 catalog in five commands

```
# 1. Pull 69 of the 207 layers of South Tyrol's WFS (~12 min)
portolan extract wfs $BZ_WFS ./bz --license CC0-1.0 \
  --layers "p_bz-Health:*,p_bz-Hydrology:Watercourses,..."

# 2. Validate against the spec
portolan check
# x PTL-PRV-001 no provider carries the 'producer' role (x69)

# 3. Who produced it, who hosts it, what it is not; 11 themes
$EDITOR .portolan/metadata.yaml

# 4. Regenerate STAC + README + AGENTS.md, re-check
portolan add --force . && portolan readme
portolan check --fix && rashid check . --schema --data
# 0 error(s), 0 warning(s) across 81 files

# 5. Push to a public bucket (Milan). Nothing else runs.
portolan push gs://south-tyrol-geodata-portolan-mirror/
rashid check . --live --live-base-url $PUBLIC_URL
# 0 error(s): range requests, CORS, sizes honoured
```

The service seeded what it could. Titles in three languages, keywords, abstracts, styles from the WMS. 69 layers to GeoParquet, spatially ordered, bbox column added; PMTiles and thumbnails alongside.

The validator failed on the machine contract, not on the prose: no producer, no license from the service. The producer's knowledge went into YAML at every level; from it came `README.md` , `collection.json` , `AGENTS.md` .

Two validators, same answer: conforms to v0.2.0, hosting contract included once pushed (range requests, CORS, sizes). The same run, hands off: an agent with the `portolan-bootstrap` skill does steps 1–5 and writes the documentation. No server was started. The result is a folder in a bucket.

[EXAMPLE CATALOG](#) [open it in the Portolan browser ↗](#) · [catalog.json on storage.googleapis.com ↗](#) · [metadata and CI on GitHub ↗](#)

Data: Provincia Autonoma di Bolzano – Alto Adige / Autonomie Provinz Bozen – Südtirol (Ripartizione 28 · ASTAT · Ripartizione 23 Salute), CC0 1.0, via data.civis.bz.it and the provincial WFS. Community mirror extracted 5 September 2026 — not an official publication; the authoritative source is the Province's geportal.

Where the metadata comes from: measured, harvested, asserted

ASSERTED BY THE PUBLISHER · `metadata.yaml`

catalog root · 1 file
license · contact · providers
attribution · keywords · citations

inherits

themes · 11 files
title · description

inherits

collections · 69 files
title · description · source_url
processing_notes · known_issues

81
metadata.yaml files, one per level

1
carries the licence, contact and providers for all 69 collections

HARVESTED · portolan extract, or an agent

WFS · CSW record · portal page
abstract · keywords · source URL · style
"TODO: Add value" if nothing was found

MEASURED · portolan add, on every run

rest-homes.parquet
extent.spatial.bbox · CRS
table:row_count · table:columns
geoparquet:geometry_type
file:size · file:checksum per asset
stac_extensions · from what was found

● ● ●
portolan add

collection.json · README.md · AGENTS.md

drafted by a tool · approved by a person · gated by rashid check in CI on every push

Story 2 · Then an agent was asked a question no single layer answers

"Which hospitals and pharmacies in South Tyrol sit inside high or very-high hydraulic hazard zones of the approved hazard-zone plans, and how many residents live in those municipalities?"

4 of 17

hospitals inside H3/H4 zones · Bolzano, Bressanone, Brunico, Ortisei

31 of 158

pharmacies inside H3/H4 zones · Merano 3, Bolzano 2, Bressanone 2, Brunico 2 ...

108 of 116

municipalities with an approved hydraulic hazard plan · 536,933 residents in the province

STATED BY THE AGENT, BECAUSE THE PUBLISHER WROTE IT DOWN

Points are locations, not service catchments. Hazard classes are planning zones, not event forecasts. Eight municipalities have no approved plan yet. Producer: Autonomous Province of Bolzano – South Tyrol · CC0-1.0 · query attached.

WHAT IT HAD TO WORK WITH

- 69 collections in 11 themes, `AGENTS.md` at every level
- the join key `ISTAT_CODE`, the CRS, the H2–H4 class strings, the caveats
- GeoParquet with bbox columns, read in place with DuckDB
- the `reading-portoLan` skill: how to navigate, query, join and report

WHAT IT DID NOT HAVE

- a server, an API, credentials, or a person in the loop

A screening, not a risk assessment. The same catalog answers the next question tomorrow.

What the agent actually did with the metadata

1

Open catalog.json

follow the explicit child links: 11 theme catalogs, 69 collections; the linked graph is the publication boundary

2

Read AGENTS.md

join key ISTAT_CODE , CRS, the H2-H4 class strings, what the data is not for

3

Inspect each collection

picks five of the 69: one GeoParquet asset each, with provider, license, extent, row count

4

Plan a DuckDB query

the model plans and explains; DuckDB computes, reading only the row groups it needs

5

Report

method, exact URLs and query, producer, license, limitations

Every step consumed something the publisher wrote down. Nothing was guessed. That is the operational meaning of AI-ready metadata.

DEMO DESK • COFFEE BREAKS TODAY AND TOMORROW

Both stories live: publish a catalog from *your* WFS, then ask the agent *your* question. And `rashid check --live` against the hosted catalog, to see the hosting contract checked too.

The skills: the know-how that made both stories run

STORY 1 · PUBLISHING

- `portolan-bootstrap` — from a WFS, an Esri service, a portal or a folder to a documented, validated, published catalog
- `portolan-cli` — init, add, check, push, sync, conversion
- `portolan-thumbnails` — framed, checked thumbnails over a basemap
- `portolan-migrate` — bring an existing catalog up to the spec without rebuilding it
- `git-backed-catalog` · `sourcecoop` · `register-catalog` — versioned metadata, hosting, the registry pull request

STORY 2 · CONSUMING

- `reading-portolan` — navigate the STAC tree, query GeoParquet with DuckDB, read COGs, join across collections, report with sources
- `portolan-consume` — explore and visualise a catalog: PMTiles, MapLibre, deck.gl
- `report-catalog-issue` — a wrong licence or a broken link becomes an issue on the registry, from the agent

Plain `SKILL.md` files in the open **Agent Skills** format: the same know-how works in Claude Code, Gemini CLI and other agents. Ten skills, one repository, Apache-2.0. **The spec says what a good catalog is; the skills encode how to build and use one.**

Lesson 1 of 3 · Discovery is the solved half. Evaluability is the gap.

WHAT A GOOD STAC CATALOG ALREADY GIVES YOU

Found. Crawled. Opened.

WHAT PEOPLE AND AGENTS STILL LACK

Is this dataset *for* my question? Where is it weak? How wrong can it be?

ONE SENTENCE FROM A REAL CATALOG (FIELDS OF THE WORLD)

A field polygon "is a *remote-sensing field unit*, **not** a cadastral/legal parcel. This is not a land-tenure product."

No accuracy figure prevents that misuse. One sentence of scope does.

Source: Fields of the World, global field-boundary dataset documentation (fieldsoftheworld, CC BY 4.0), as quoted in the Portolan documentation best practices.

AI-ready = documented limitations, quantified uncertainty, named biases — **as data where possible** (a confidence column, a quality mask), not buried in prose. Expertise that otherwise lives only in the heads of domain experts.

WHERE A LIMITATION LIVES IN A PORTOLAN CATALOG

`known_issues:`

in `metadata.yaml`, written once by the publisher



`## Known Issues`

in `README.md`, for people



`## Data quality & usage notes`

in `AGENTS.md`, for agents



`confidence` column

as data, where the limitation is per row

Lesson 2 of 3 · Two audiences, two files, one subject

ONE COLLECTION · THREE ENTRANCES



THE SAME PUBLISHED ASSETS

GeoParquet PMTiles · styles · thumbnail · original source

README.md helps a person decide whether to trust the data. AGENTS.md helps an agent that already committed to using it get the query right the first time. Both generated from one source.

Lesson 3 of 3 · Be strict, and let a machine hold the line

WHY SO PRESCRIPTIVE

An agent can only rely on what is *required*. Every MAY in a specification is a branch every consumer has to write.

- GeoParquet without bbox columns and spatial ordering: spatial queries read the whole file
- a bucket without CORS: unusable from a browser
- an empty `providers` list: the prose looks fine, every machine fails

WHAT THE MACHINE HOLDS · SPEC v0.2.0

128

normative
requirements

105

checked
deterministically by
rashid

23

stay with the
publisher's judgment

MUST → failure, SHOULD → warning, JSON out. A city's catalog passes the same validator a planetary archive passes.

Conformance is a floor. "You can have perfectly FAIR, utterly useless data." Above pass/fail, a grader scores usability from A+ to F, and weights the machine contract equally with what a browser shows.

A metadata quality policy only sticks when a machine can check it.

Honestly unfinished: trust and provenance



How should **authority** be signalled cryptographically, not just declared in a providers list?



What does **trust** mean when the consumer is an agent acting unattended?



How does a **mirror** stay verifiably faithful to its source?

Today: every collection names producer, licensor, processor and host; official or mirror is derived from that; a registry of ~20 independently hosted catalogs, plain JSON in git. All declarative. The questions above map onto the OGC Testbed on Trusted Data and Systems.

Public since 2 September 2026. Early adopters wanted; breaking changes still expected. First goal: ~100 reference catalogs, then a stable v1.0 CLI.

portolan-sdi.org · [#portolan](#) on the Cloud-Native Geospatial Forum Slack · community call Fridays 10:00 CET · demo desk at the coffee breaks

Questions? Everything is open and one pull request away.

PROJECT LINKS

CLICK ANY CARD

WEBSITE

portolan-sdi.org



ALL PROJECTS

github.com/portolan-sdi



SPECIFICATION

portolan-spec



PUBLISHER CLI

portolan-cli



VALIDATOR

rashid



AGENT SKILLS

portolan-skills



PUBLIC REGISTRY

portolan-registry



CATALOG BROWSER

browser.portolan-sdi.org



THIS DECK




cayetanobv.github.io

ogc-metadata-summit-bolzano/#/17

SCAN · CLICK · SHARE

Appendix 1 • Why "Portolan"



Portolan chart of the Mediterranean and the Black Sea, Battista Agnese, 1550. Bibliothèque nationale de France, via Wikimedia Commons (public domain).

Portolan charts were among the earliest **practical** navigational maps: not theoretical pictures of the world but working tools, built from sailors' observations, continuously refined, and shared across ports and nations.

- **Decentralised** — made and improved by many hands
- **Interoperable** — readable wherever a ship went
- **Critical infrastructure** — for trade, for safety, for sovereignty

The project takes its name from that tradition. Practical rather than academic; built from data in open formats rather than drawings; designed for sharing across organisations, clouds and borders.

Appendix 2 · Your ISO 19115 / INSPIRE records are inputs, not casualties

ISO IN

`portolan extract wfs` reads the ISO 19139 record over CSW and seeds title, abstract, keywords, license, contact, lineage, dates.

ISO ALONGSIDE

The record travels with the data as an asset with `roles`:
`["metadata","iso-19115"]` and a typed `describedby` link to the catalogue that owns it.

GEODCAT-AP OUT

Generated from the same source as a sibling representation, so open-data portals can harvest it.
Roadmap, not shipped.

The same pattern swisstopo and the CEOS STAC best practices already follow. There is still no ISO 19115 STAC extension — an open question for this room.

Appendix 2 • How an ISO record rides in a Portolan collection

```
"assets": {
  "metadata": {
    "href": "./iso19115.xml",
    "type": "application/vnd.iso.19139+xml",
    "title": "ISO 19115 record, BRK Bestuurlijke Gebieden (Nationaal Georegister)",
    "roles": ["metadata", "iso-19115"],
    "file:size": 40942, "file:checksum": "1220..."
  }
},
"links": [
  {"rel": "describedby", "href": "./README.md", "type": "text/markdown"},
  {"rel": "describedby",
    "href": "https://www.nationaalgeoregister.nl/geonetwork/srv/api/records/208bc283-.../formatters/xml",
    "type": "application/vnd.iso.19139+xml", "title": "ISO 19115 record – authoritative"},
  {"rel": "via", "href": "https://www.pdok.nl/...", "type": "text/html", "title": "Original source (PDOK)"}
]
```

Asset = pinned copy with checksum and mirror date. Link = live authoritative record. If they disagree, the link wins and the asset tells you how old the copy is.